

CIFAR
CANADIAN
INSTITUTE
FOR
ADVANCED
RESEARCH



Deep Learning Summer School 2015

Introduction to Machine Learning

by **Pascal Vincent**



Montreal Institute for Learning Algorithms

August 4, 2015

Université
de Montréal

Département d'informatique et de recherche
opérationnelle

What is machine learning ?

Historical perspective


- Born from the ambitious goal of **Artificial Intelligence**



- Founding project:

The **Perceptron** (Frank Rosenblatt 1957)

First artificial neuron **learning** form examples



- Two historically opposed approaches to AI:

Neuroscience inspired:

→ neural nets **learning from examples** for artificial perception

Classical symbolic AI:

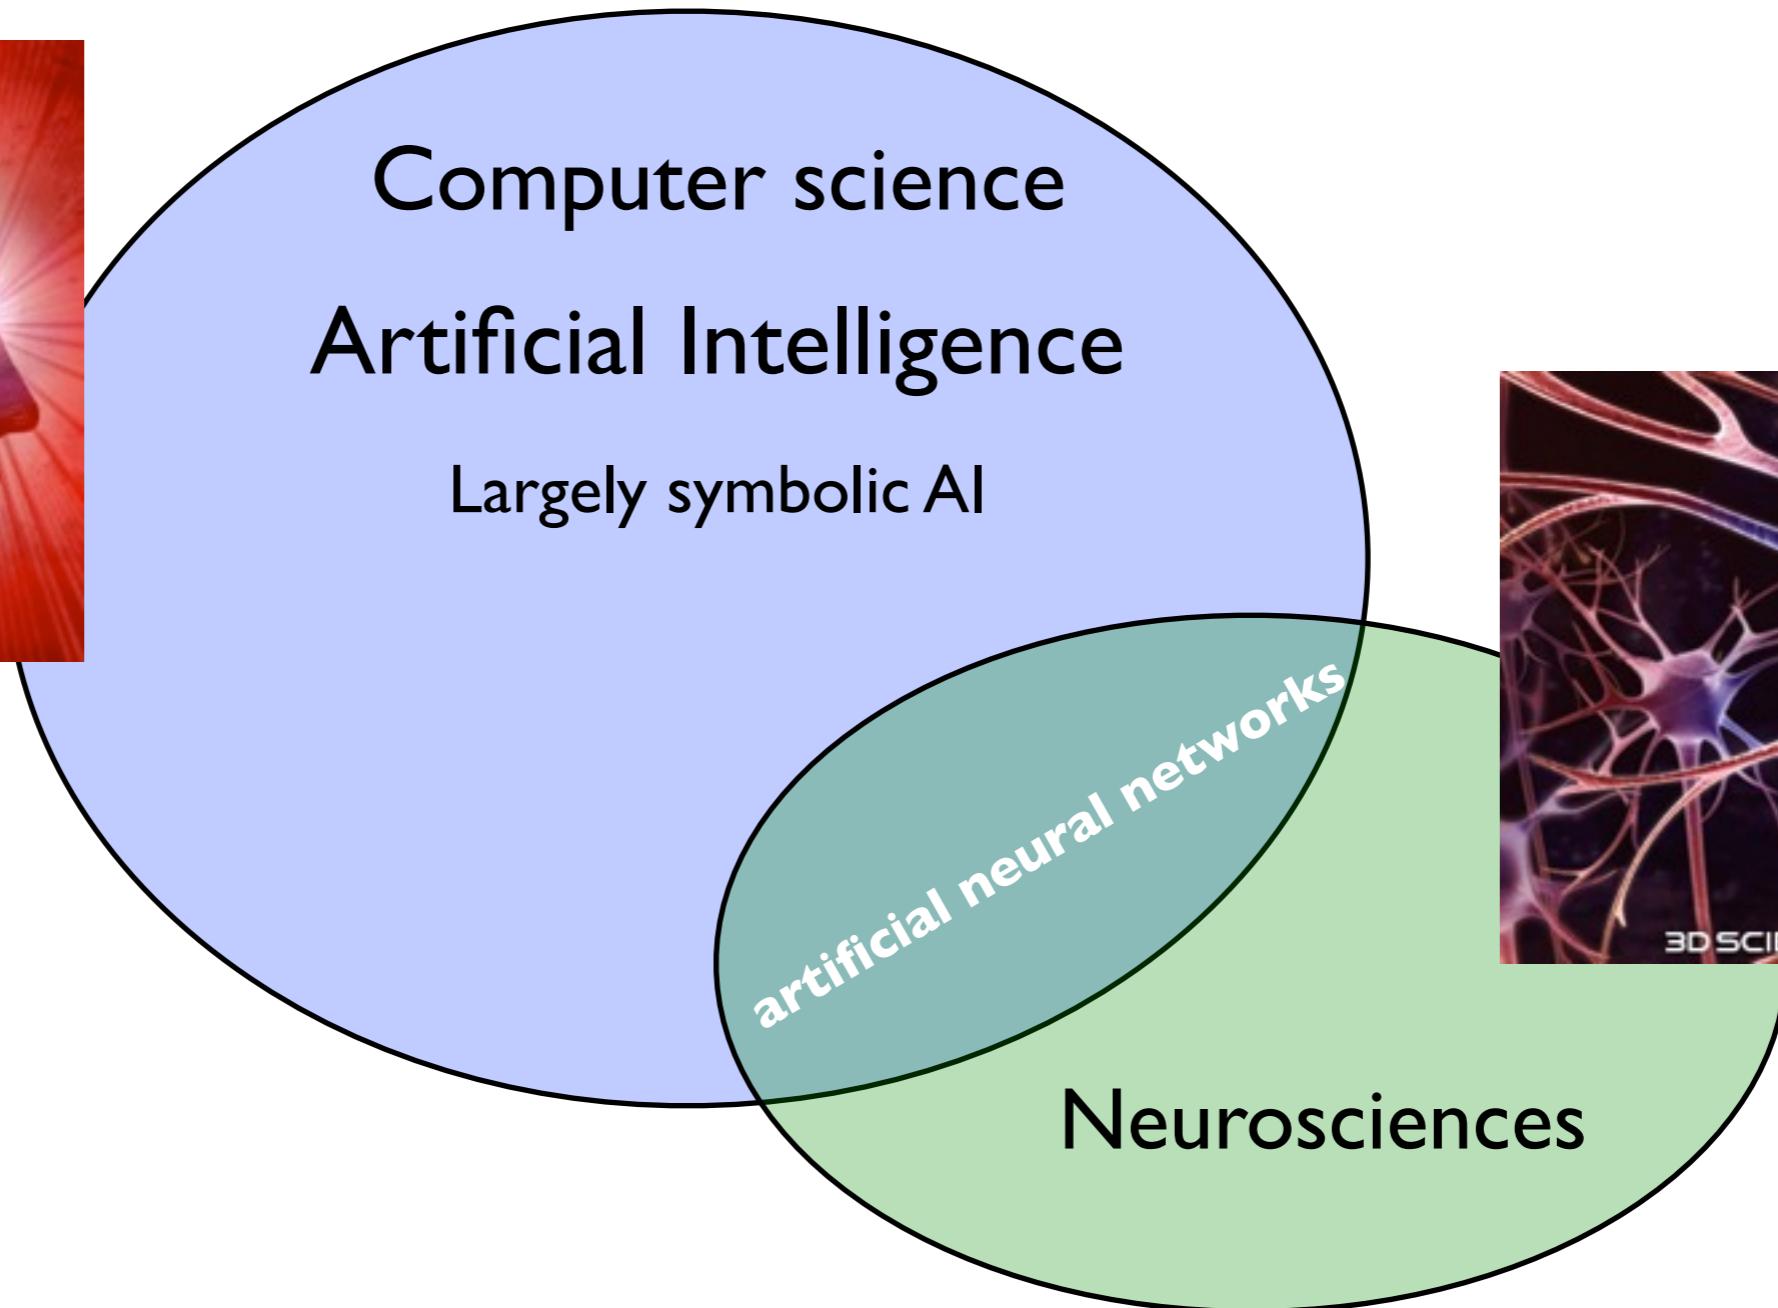
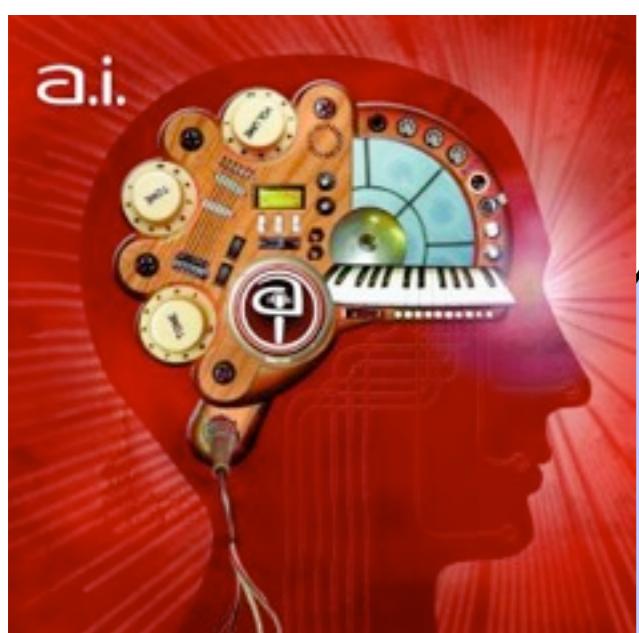
Primacy of **logical reasoning** capabilities

→ No learning (humans coding rules)
→ poor handling of uncertainty


Got eventually fixed (Bayes Nets...)

Learning and probabilistic models largely won → machine learning

Artificial Intelligence in the 60s



Current view of ML founding disciplines



What is machine-learning?



A (hypnotized) user's perspective

A scientific (**witchcraft**) field that

- researches fundamental principles (**potions**)
- and **develops** magical algorithms (**spells to invoke**)
- capable of **leveraging collected data** to (**automagically**)
produce accurate predictive functions
applicable to similar data (**in the future!**)

(may also yield informative *descriptive* functions of data)

The key ingredient of machine learning is...



Data!

- Collected from nature... or industrial processes.
- Comes stored in many forms (and formats...), structured, unstructured, occasionally clean, **usually messy**, ...
- In ML we like to view data as a **list of examples**
 (or we'll turn it into one)
 - ➡ ideally **many examples** of the **same nature**.
 - ➡ preferably with each example a **vector of numbers**
(or we'll first turn it into one!)

D_n

Training data set (training set)



New test point:



→ ?

Input dimensionality:

d

| inputs: X (input feature vector) | targets: Y (label) |
|--|-------------------------|
| X_1 (3.5, -2, ..., 127, 0, ...) | +1 |
| X_n (-9.2, 32, ..., 24, 1, ...) | -1 |
| $X_{n,2}$ (6.8, 54, ..., 17, -3, ...) | +1 |

$$X = (5.7, -27, \dots, 64, 0, \dots) \xrightarrow{f_\theta} +1$$

$$\mathbf{x} \in \mathbb{R}^d$$

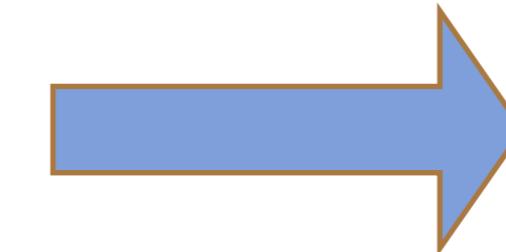
Importance of the Problem dimensions

- ⇒ Détermines which learning algorithms will be practically applicable (based on their algorithmic complexity and memory requirements).
 - Number of examples: **n**
(sometimes several millions)
 - Input dimensionality: **d**
number of input features characterizing each example
(often 100 to 1000, sometimes 10000 or much more)
 - Target dimensionality ex. number of classes **m**
(often small, sometimes huge)
- Data suitable for ML will often be organized as a matrix: **$n \times (d+1)$** ou **$n \times (d+m)$**

Turning ~~messy~~ data into a nice list of examples



data-plumbing

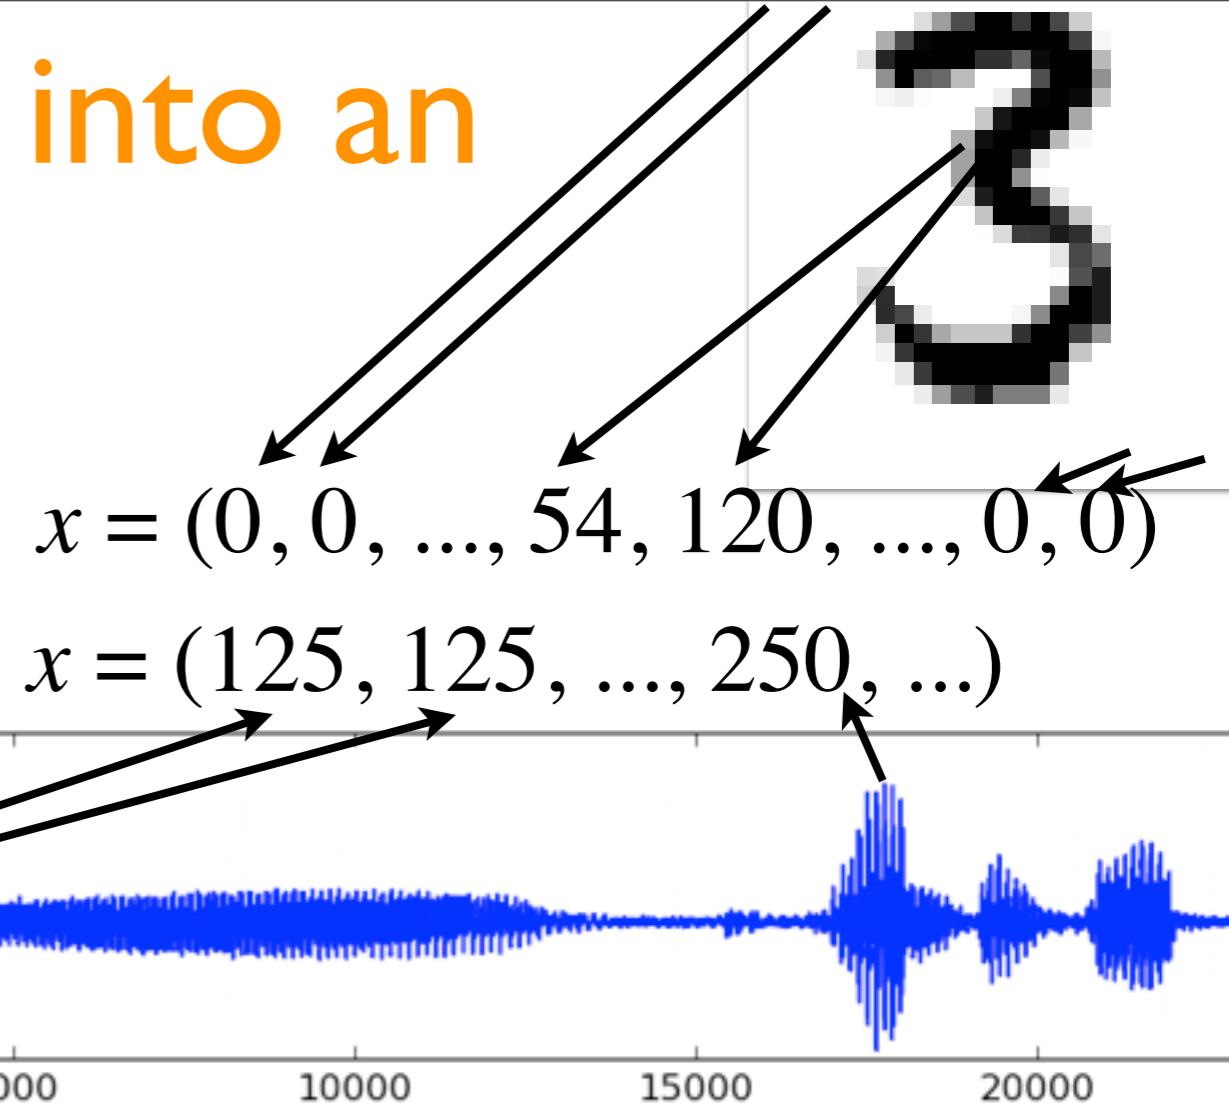


Key questions to decide what «examples» should be:

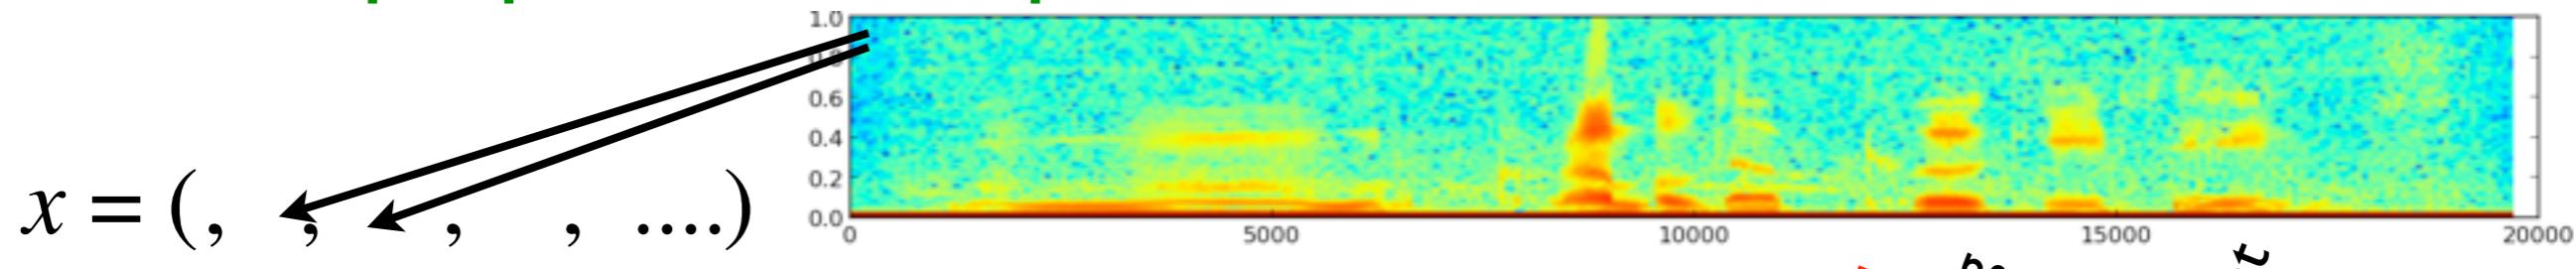
- **input:** What is all the (potentially relevant) **information** I will have at my disposal about a case when I will have to make a prediction about it? (at test time)
- **target:** what I want to predict: Can I get my hands on many such examples that are actually labeled with prediction targets?

Turning an example into an input vector $\mathbf{x} \in \mathbb{R}^d$

Raw input representation:



OR some preprocessed representation:

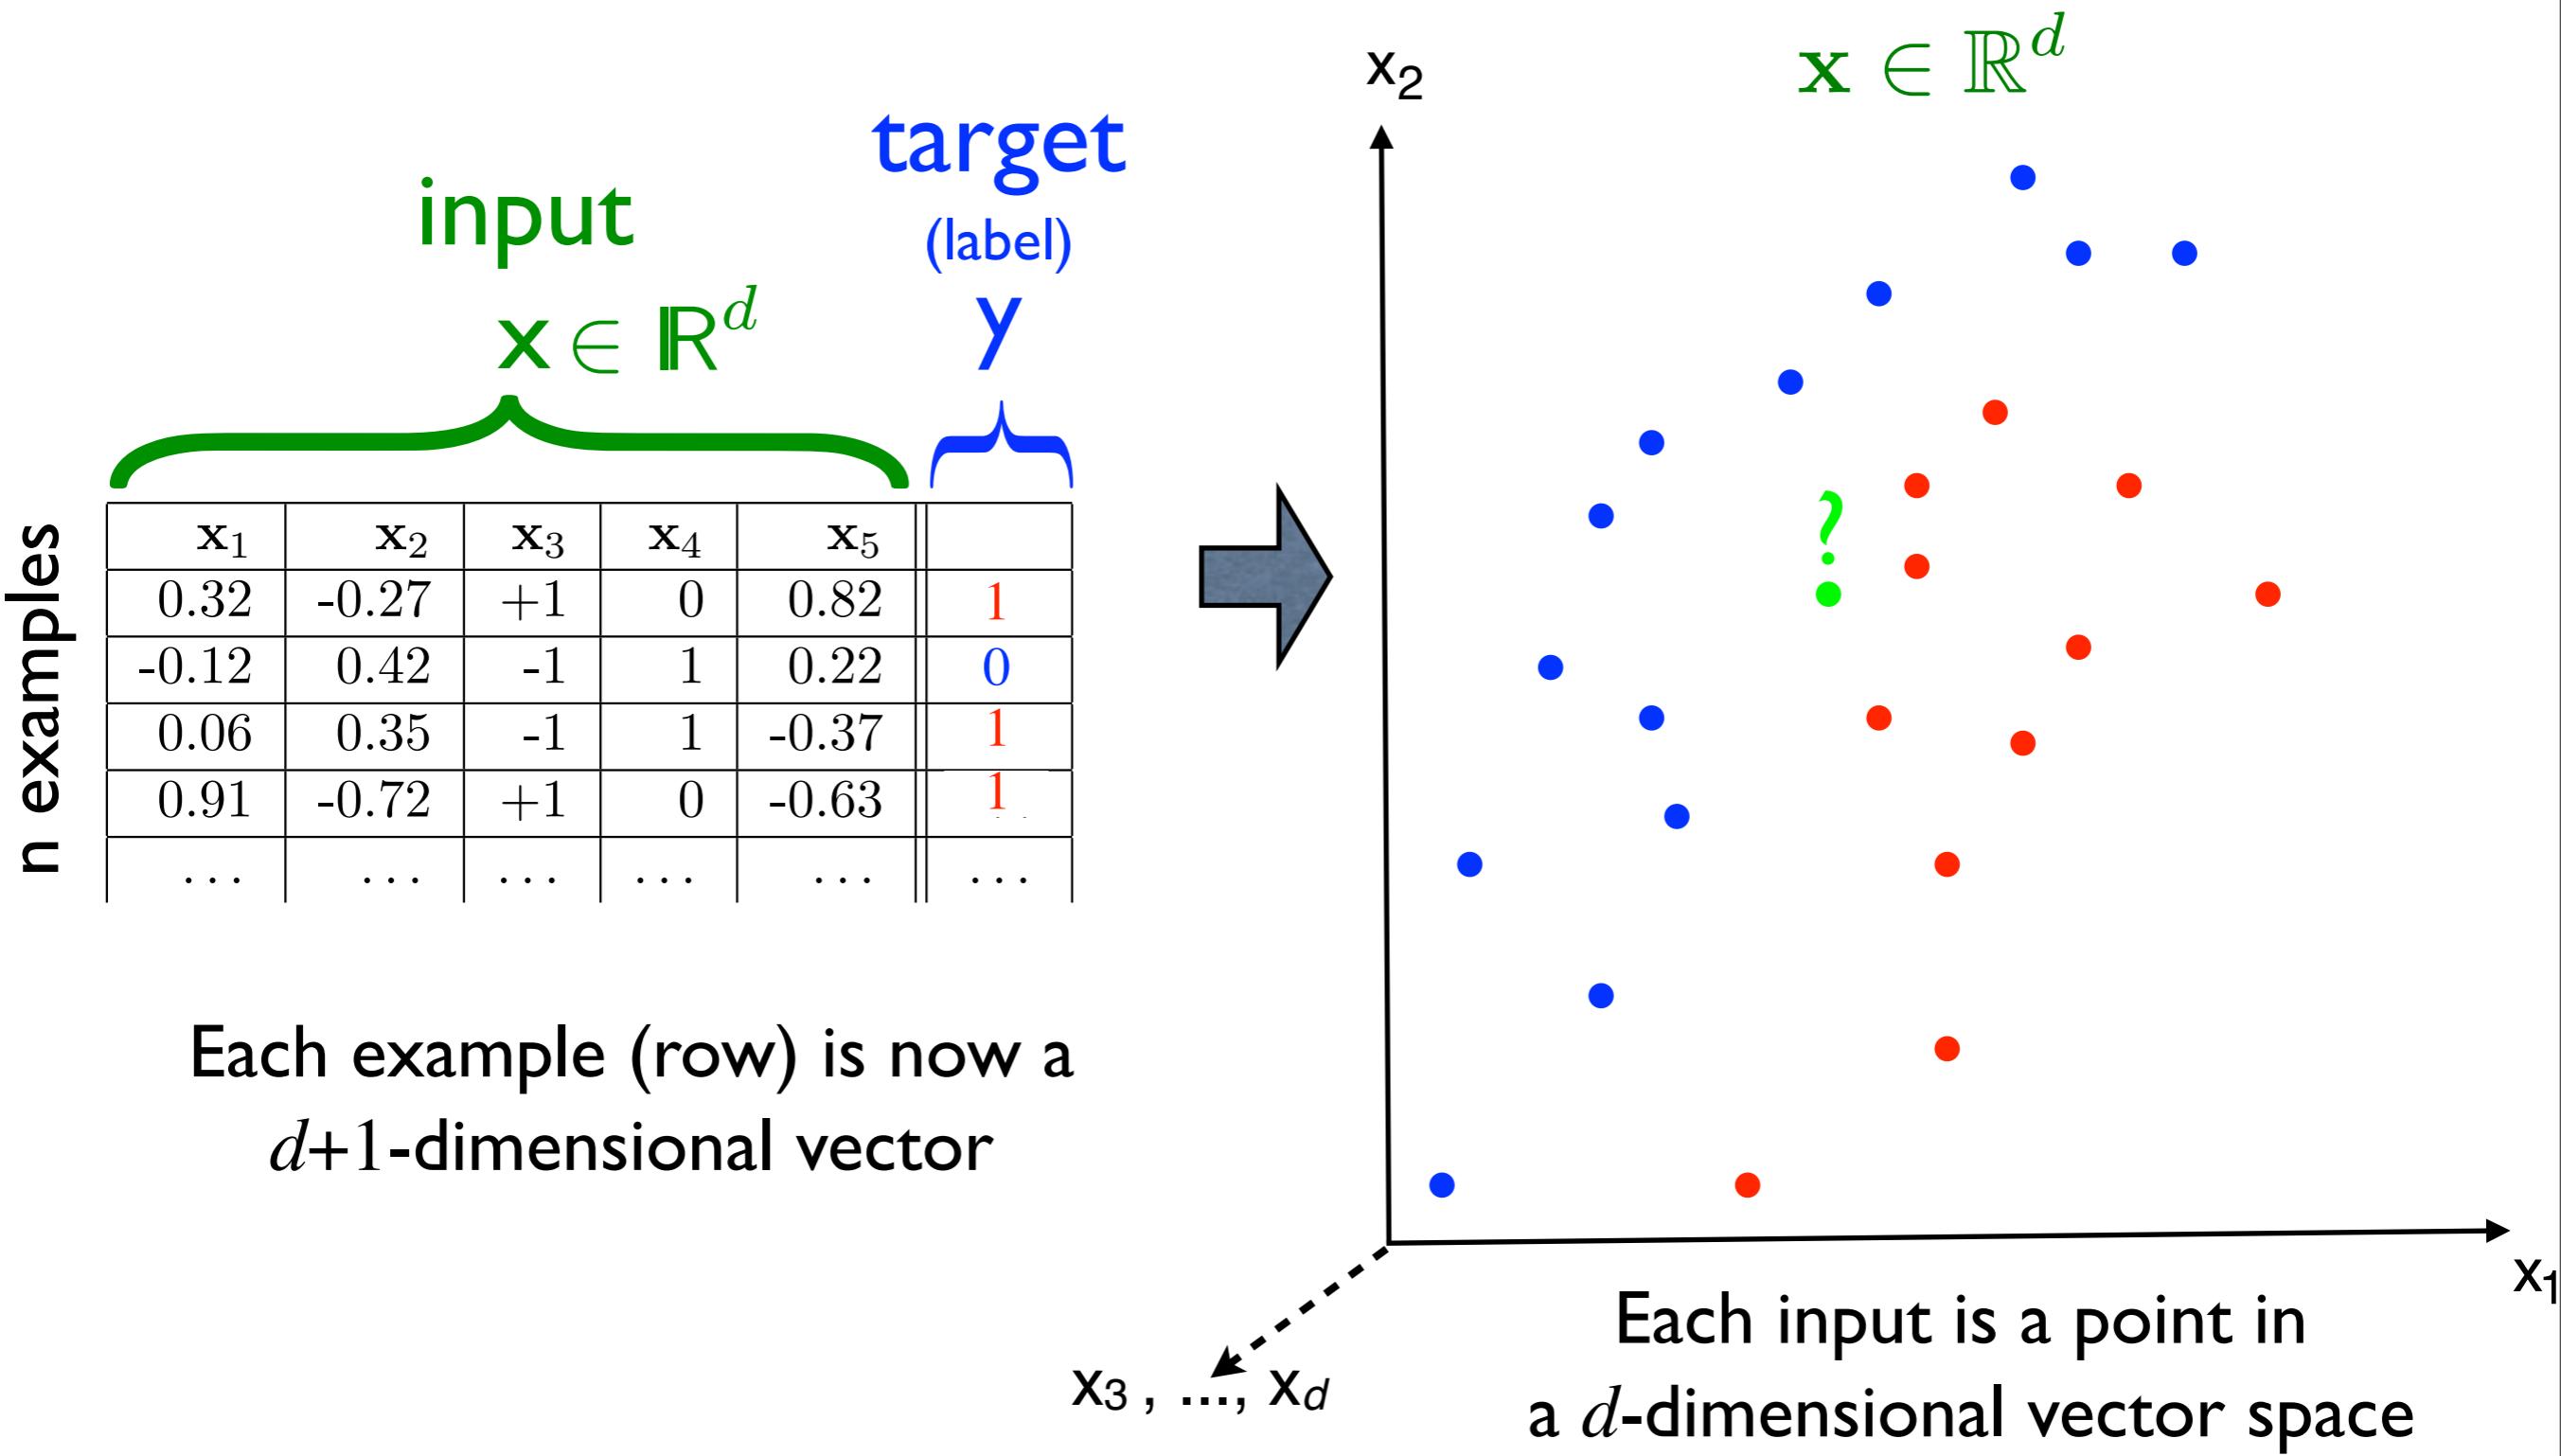


Bag of words for «The cat jumped»: $\mathbf{x} = (\dots 0 \dots, 0, 1, \dots 0 \dots, 1, 0, 0, \dots, 0, 0, 1, 0, \dots 0 \dots)$

OR vector of hand-engineered features:
ex: Histograms of Oriented Gradients

$\mathbf{x} = (\text{feature } 1, \dots, \text{feature } d)$

Dataset imagined as a point cloud in a *high-dimensional* vector space




Ex: nearest-neighbor classifier

Algorithm:

For test point x :

- Find **nearest neighbor** of x among the training set according to some distance measure (eg: Euclidean distance).
- Predict that x has the same class as this nearest neighbor.



Machine learning tasks (problem types)

Supervised learning = predict a target y from input x

(and semi-supervised learning)

- y represents a category or “class”

⇒ classification

binary : $y \in \{-1, +1\}$ or $y \in \{0, 1\}$

multiclass : $y \in \{1, m\}$ or $y \in \{0, m - 1\}$

- y is a real-value number

⇒ regression

$y \in \mathbb{R}$ or $y \in \mathbb{R}^m$

Predictive models

Unsupervised learning: no explicit prediction target y

- model the probability distribution of x

⇒ density estimation

- discover underlying structure in data

⇒ clustering

⇒ dimensionality reduction

⇒ (unsupervised) representation learning


Descriptive modeling

Reinforcement learning: taking good sequential decisions to maximize a reward in an environment influenced by your decisions.

Learning phases

- **Training:** we **learn** a predictive function f_θ by optimizing it so that it predicts well on **the training set**.
 - **Use for prediction:** we can then use f_θ on new (test) inputs that were not part of the training set.
- ⇒ The GOAL of learning is *NOT* to learn perfectly (*memorize*) the training set.
- ⇒ What's important is the ability for the predictor to **generalize** well on new (future) cases.

Ex: 1D regression




1. Collect training data
2. Learn a function (predictor)
input → target
3. Use learned function
on new inputs

Original slide by Olivier Delalleau

Supervised task:

predict y from x



Training set D_n

Learn a function f_θ that will minimize prediction errors as measured by cost (loss) L .

loss function

$$L(f_\theta(\mathbf{x}), \mathbf{y})$$

output

$$f_\theta(\mathbf{x})$$

$$f_\theta$$

: parameters

| | | | | |
|-------|------|----|---|------|
| -0.12 | 0.42 | -1 | 1 | 0.22 |
|-------|------|----|---|------|

34

target y

input \mathbf{x}

A machine learning algorithm
usually corresponds to a combination of
the following 3 elements:
(either explicitly specified or implicit)

- ✓ the choice of a specific **function family**: F
(often a parameterized family)
- ✓ a **way to evaluate the quality** of a function $f \in F$
(typically using a **cost** (or **loss**) function L
measuring how wrongly f predicts)
- ✓ a **way to search for the «best»** function $f \in F$
(typically an optimization of function parameters to
minimize the overall loss over the training set).

Evaluating the quality of a function $f \in F$

and

Searching for the «best» function $f \in F$

Evaluating a predictor $f(x)$

The performance of a predictor is often **evaluated using several different evaluation metrics**:

- Evaluations of **true quantities of interest** (\$ saved, #lifes saved, ...) when using predictor inside a more complicated system.
- «Standard» evaluation metrics in a specific field (e.g. BLEU (Bilingual Evaluation Understudy) scores in translation)
- Misclassification error rate for a classifier (or precision and recall, or F-score, ...).
- **The loss actually being optimized** by the ML algorithm (often different from all the above...)

Standard loss-functions

- **For a density estimation task:** $f : \mathbb{R}^d \rightarrow \mathbb{R}^+$ a proper probability mass or density function
negative log likelihood loss: $L(f(x)) = -\log f(x)$
- **For a regression task:** $f : \mathbb{R}^d \rightarrow \mathbb{R}$
squared error loss: $L(f(x), y) = (f(x) - y)^2$
- **For a classification task:** $f : \mathbb{R}^d \rightarrow \{0, \dots, m-1\}$
misclassification error loss: $L(f(x), y) = I_{\{f(x) \neq y\}}$

Surrogate loss-functions

- For a classification task: $f : \mathbb{R}^d \rightarrow \{0, \dots, m-1\}$
 misclassification error loss: $L(f(x), y) = I_{\{f(x) \neq y\}}$

Problem: it is hard to optimize the misclassification loss directly
 (gradient is 0 everywhere. NP-hard with a linear classifier) Must use a surrogate loss:

| | Binary classifier | Multiclass classifier |
|------------------------------|---|--|
| Probabilistic classifier | <p>Outputs probability of class 1 $g(x) \approx P(y=1 x)$ Probability for class 0 is $1-g(x)$</p> <p><u>Binary cross-entropy loss:</u> $L(g(x), y) = -(y \log(g(x)) + (1-y) \log(1-g(x)))$</p> <p>Decision function: $f(x) = I_{g(x)>0.5}$</p> | <p>Outputs a vector of probabilities: $g(x) \approx (P(y=0 x), \dots, P(y=m-1 x))$</p> <p><u>Negated conditional log likelihood loss</u> $L(g(x), y) = -\log g(x)_y$</p> <p>Decision function: $f(x) = \text{argmax}(g(x))$</p> |
| Non-probabilistic classifier | <p>Outputs a «score» $g(x)$ for class 1. score for the other class is $-g(x)$</p> <p><u>Hinge loss:</u> $L(g(x), t) = \max(0, 1-tg(x))$ where $t=2y-1$</p> <p>Decision function: $f(x) = I_{g(x)>0}$</p> | <p>Outputs a vector $g(x)$ of real-valued scores for the m classes.</p> <p><u>Multiclass margin loss</u> $L(g(x), y) = \max(0, 1 + \max_{k \neq y} (g(x)_k - g(x)_y))$</p> <p>Decision function: $f(x) = \text{argmax}(g(x))$</p> |

Expected risk v.s. Empirical risk

Examples (\mathbf{x}, \mathbf{y}) are supposed drawn i.i.d. from an **unknown true distribution** $p(\mathbf{x}, \mathbf{y})$ (from nature or industrial process)

- **Generalization error** = Expected risk (or just «Risk»)
«how poorly we will do on average on the infinity of future examples from that unknown distribution»

$$R(f) = \mathbb{E}_{p(\mathbf{x}, \mathbf{y})}[L(f(\mathbf{x}), \mathbf{y})]$$

- **Empirical risk** = average loss on a finite dataset
«how poorly we're doing on average on this finite dataset»

$$\hat{R}(f, D) = \frac{1}{|D|} \sum_{(\mathbf{x}, \mathbf{y}) \in D} L(f(\mathbf{x}), \mathbf{y})$$

where $|D|$ is the number of examples in D

Empirical risk minimization

Examples (x,y) are supposed drawn i.i.d. from an unknown true distribution $p(x,y)$ (nature or industrial process)

- We'd **love** to find a predictor that **minimizes the generalization error** (the expected risk)
- But **can't even compute it!** (expectation over unknown distribution)
- Instead: **Empirical risk minimization principle**
«Find predictor that minimizes average loss over a trainset»

$$\hat{f}(D_{\text{train}}) = \underset{f \in F}{\operatorname{argmin}} \hat{R}(f, D_{\text{train}})$$

This is the training phase in ML

Evaluating the generalization error

- ▶ We can't compute expected risk $R(f)$
- ▶ But $\hat{R}(f, D)$ is a good estimate of $R(f)$ provided:
 - D was not used to find/choose f otherwise estimate is biased \Rightarrow can't be the training set!
 - D is large enough (otherwise estimate is too noisy); drawn from p

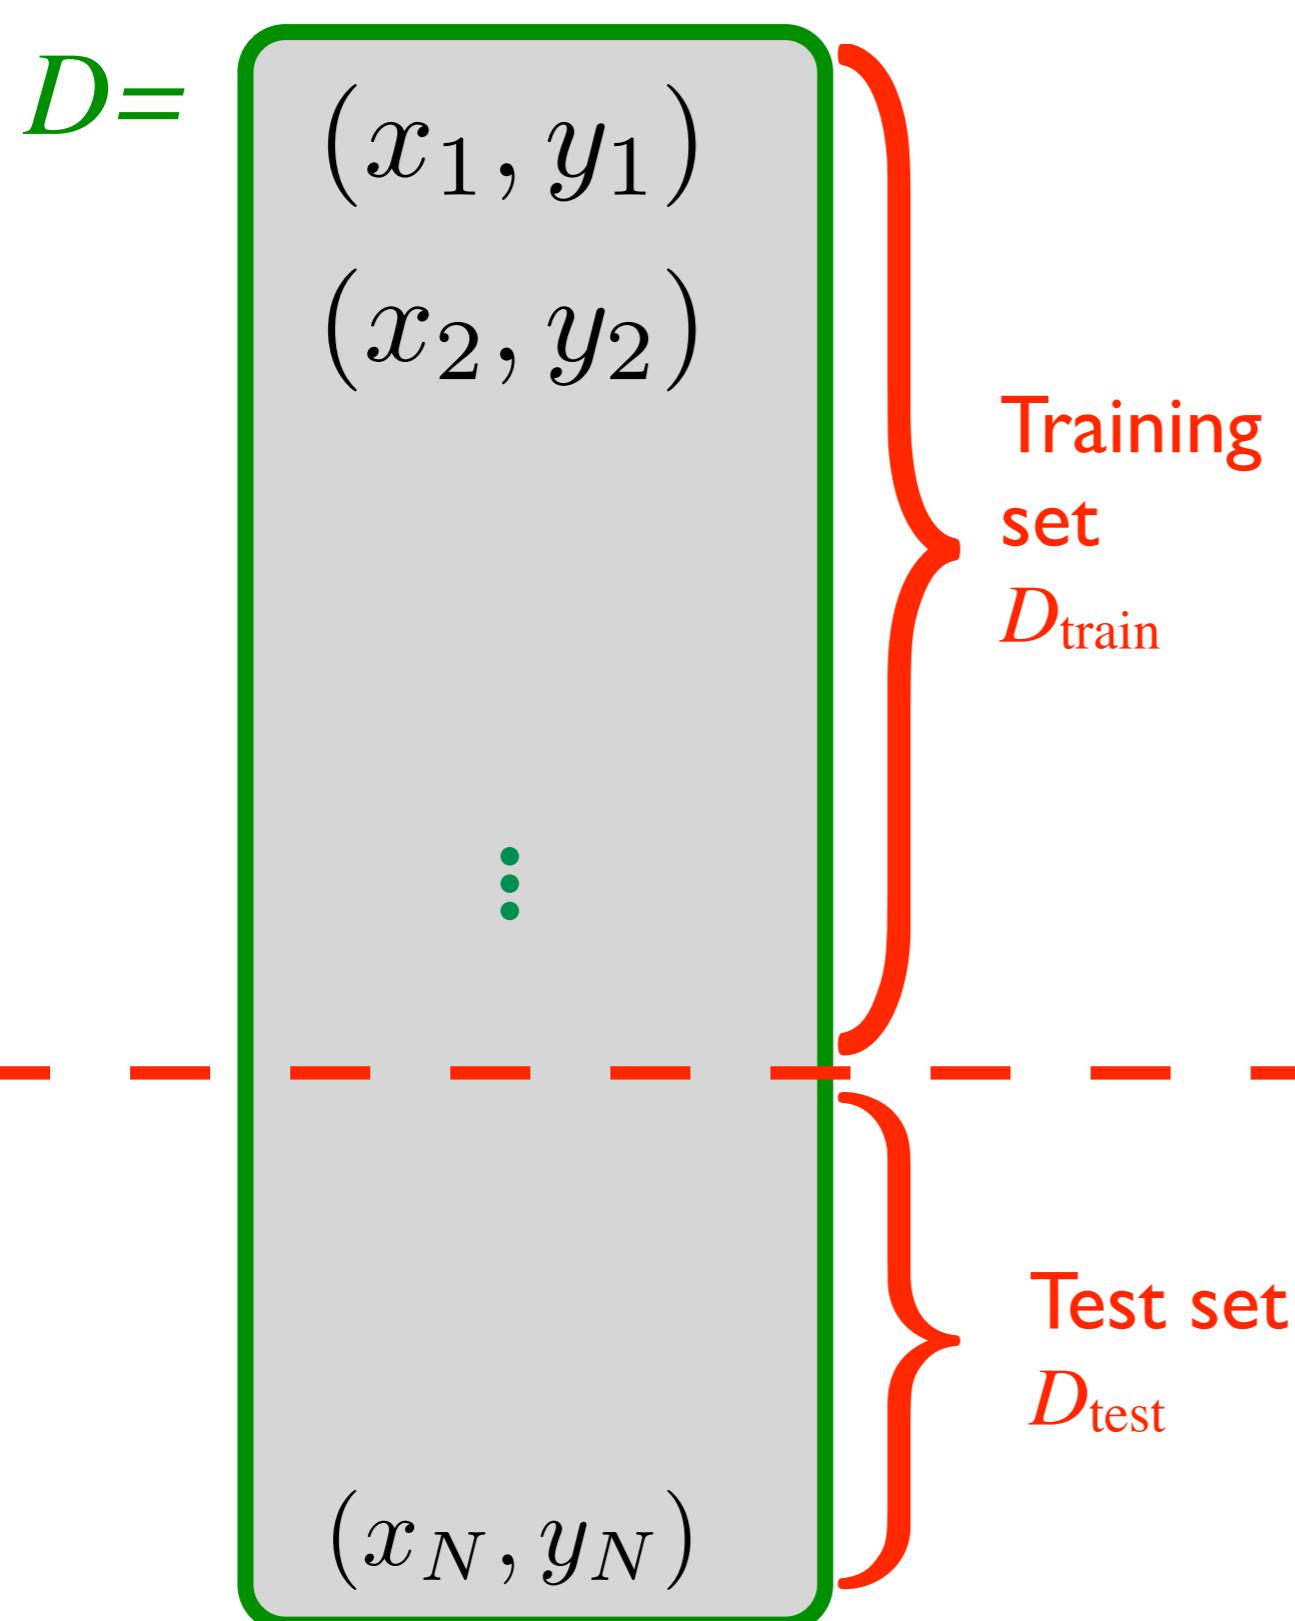
→ Must keep a separate test-set $D_{\text{test}} \neq D_{\text{train}}$ to properly estimate generalization error of $\hat{f}(D_{\text{train}})$:

$$R(\hat{f}(D_{\text{train}})) \approx \hat{R}(\hat{f}(D_{\text{train}}), D_{\text{test}})$$

generalization average error on
error test-set (never used for training)

This is the test phase in ML

Simple train/test procedure



- Provided large enough dataset D drawn from $p(x,y)$
- Make sure examples are in random order.
- Split dataset in **two**: D_{train} and D_{test}
- Use D_{train} to choose/optimize/find best predictor $f = \hat{f}(D_{\text{train}})$
- Use D_{test} to evaluate generalization performance of predictor f .

Model selection


Choosing a specific
function family F

Ex. of parameterized function families

$F_{\text{polynomial } p}$

Polynomial predictor (of degree p):

$$f(x) = b + a_1x + a_2x^2 + a_3x^3 + \dots + a_p x^p$$



F_{linear}


Linear (affine) predictor:
«linear regression»

$$\theta = \{w \in \mathbb{R}^d, b \in \mathbb{R}\}$$

$$f_\theta(x) = wx + b \quad (\text{in 1 dimension})$$

$$f_\theta(x) = w^T x + b \quad (\text{in } d \text{ dimensions})$$

Q:
pre
test



F_{const}

Constant predictor: $f_\theta(x) = b$
where $\theta = \{b\}$
(always predict the same value or class!)

Capacity of a learning algorithm

- Choosing a specific Machine Learning algorithm means choosing a specific function family F .
- How «big, rich, flexible, expressive, complex» that family is, defines what is informally called the «capacity» of the ML algorithm.
Ex: $\text{capacity}(F_{\text{polynomial } 3}) > \text{capacity}(F_{\text{linear}})$
- One can come up with several formal measures of «capacity» for a function family / learning algorithm (e.g. **VC-dimension** Vapnik–Chervonenkis)
- One rule-of-thumb estimate, is the **number of adaptable parameters**: i.e. how many scalar values are contained in θ .
Notable exception: chaining many linear mappings is still a linear mapping!

Effective capacity, and capacity-control hyper-parameters

The «effective» capacity of a ML algo is controlled by:

- Choice of ML algo, which determines big family F
- Hyper-parameters that further specify F
e.g.: degree p of a polynomial predictor; Kernel choice in SVMs;
#of layers and neurons in a neural network
- Hyper-parameters of «regularization» schemes
e.g. constraint on the norm of the weights w
(\Rightarrow ridge-regression; L_2 weight decay in neural nets);
Bayesian prior on parameters; noise injection (dropout); ...
- Hyper-parameters that control early-stopping of the iterative search/optimization procedure.
(\Rightarrow won't explore as far from the initial starting point)

Popular classifiers


their parameters and hyper-parameters

| Algo | Capacity-control hyperparameters | Learned parameters |
|---|---|---|
| logistic regression (L ₂ regularized) | strength of L ₂ regularizer | w,b |
| linear SVM | C | w,b |
| kernel SVM | C; kernel choice & params (σ for RBF; degree for polynomial) | support vector weights: α |
| neural network | layer sizes; early stop; ... | layer weight matrices |
| decision tree | depth | the tree (with index and threshold of variables) |
| k-nearest neighbors | k; choice of metric | memorizes trainset |

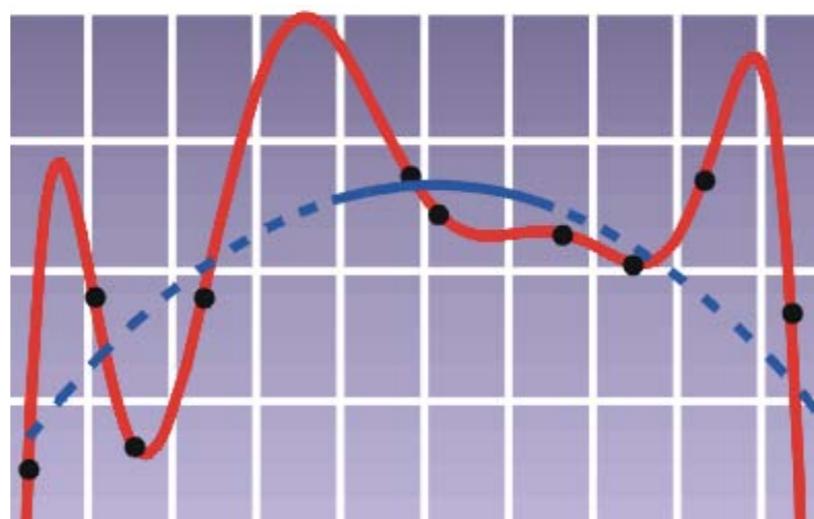
Tuning the capacity

- Capacity must be optimally tuned to ensure good generalization
- by choosing Algorithm and hyperparameters
- to avoid under-fitting and over-fitting.


Ex: 1D regression with polynomial predictor



capacity too low
→under-fitting



capacity too high
→over-fitting



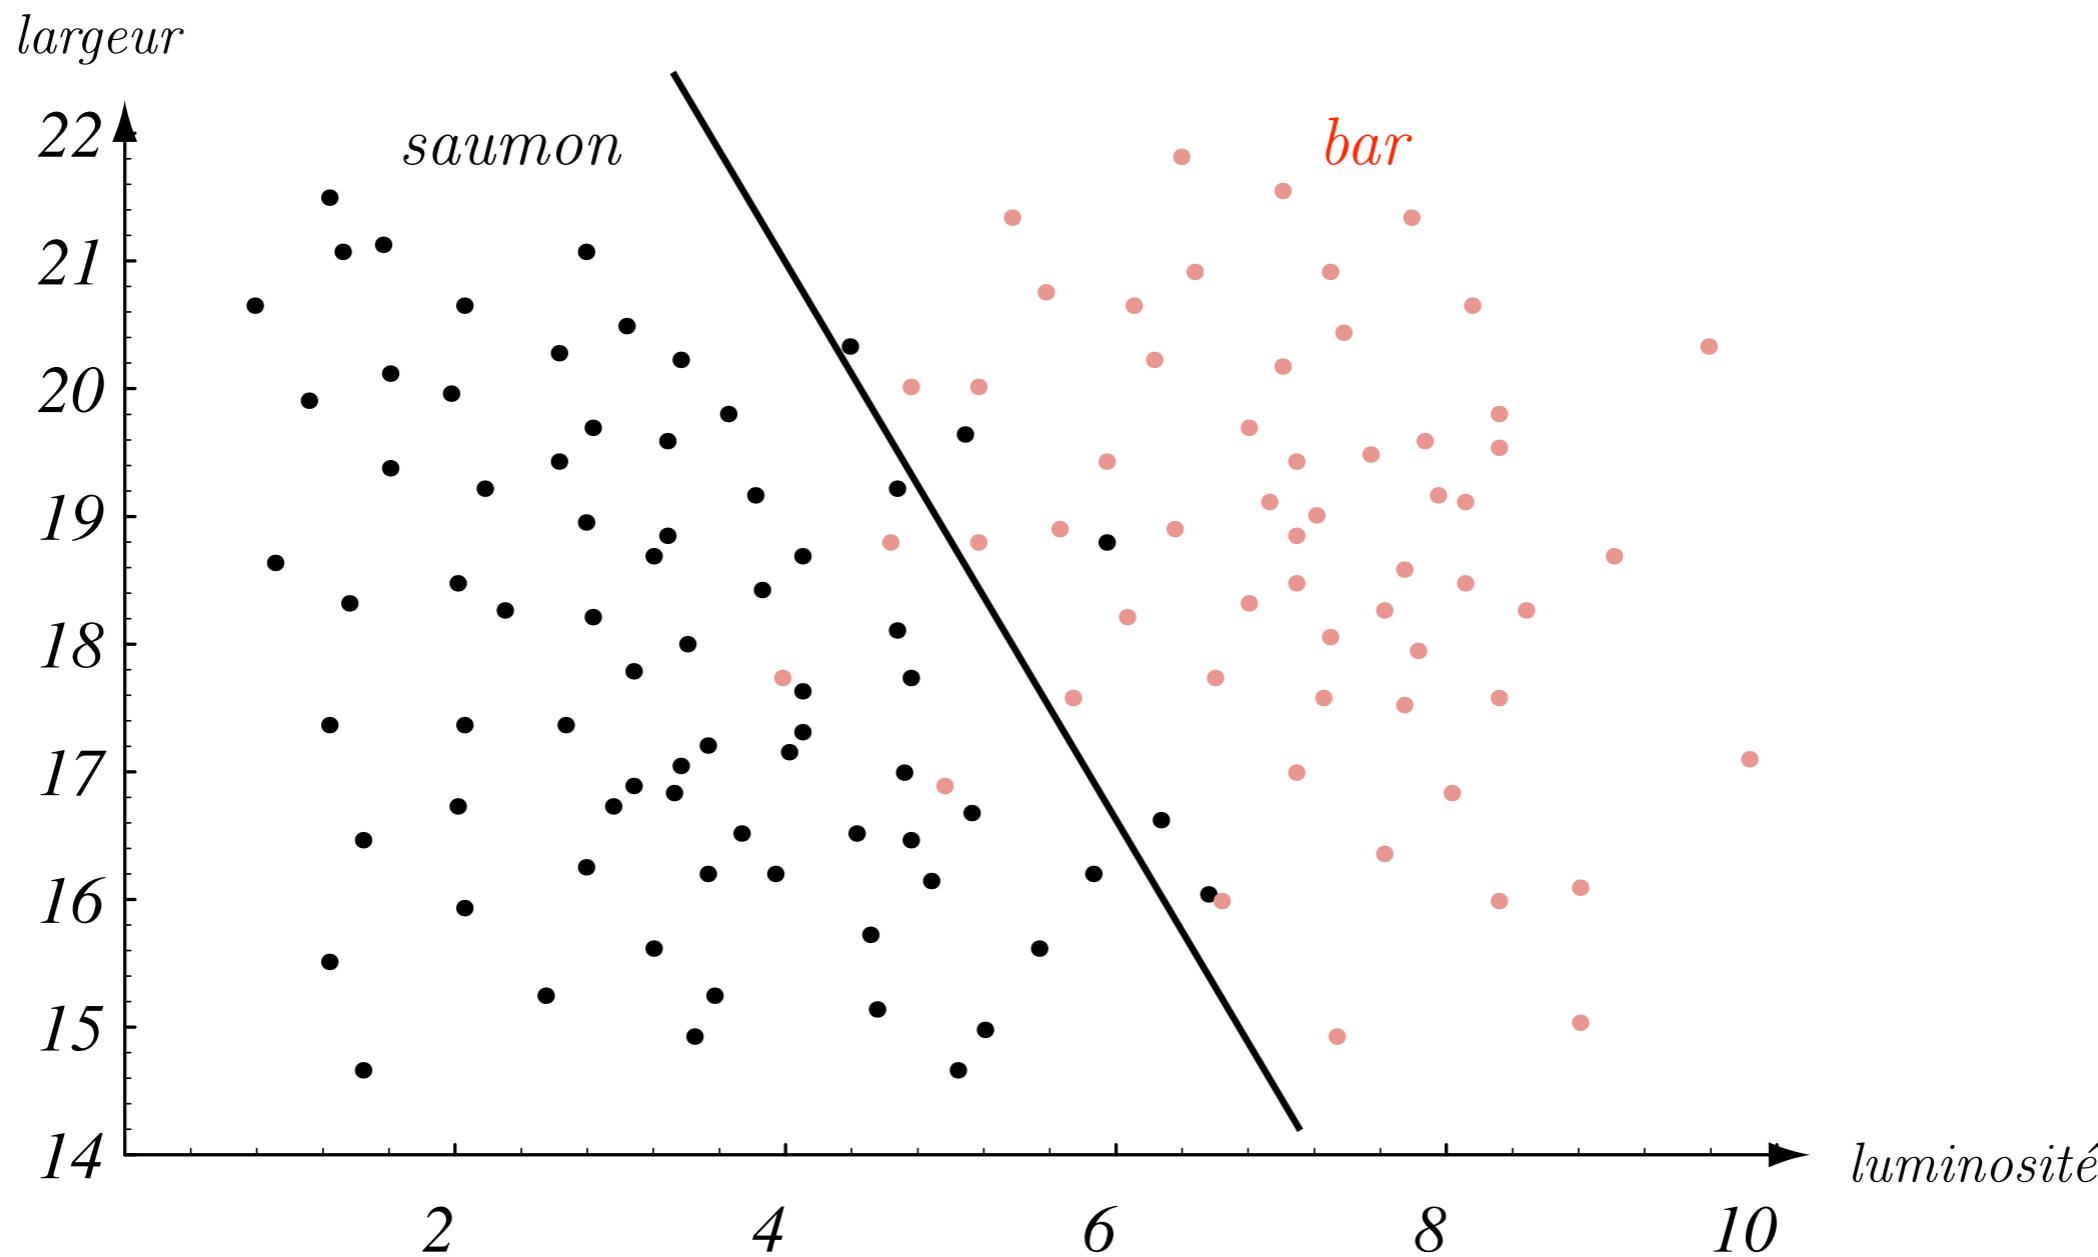
optimal capacity
→good generalisation

performance on training set is not a good estimate of generalization

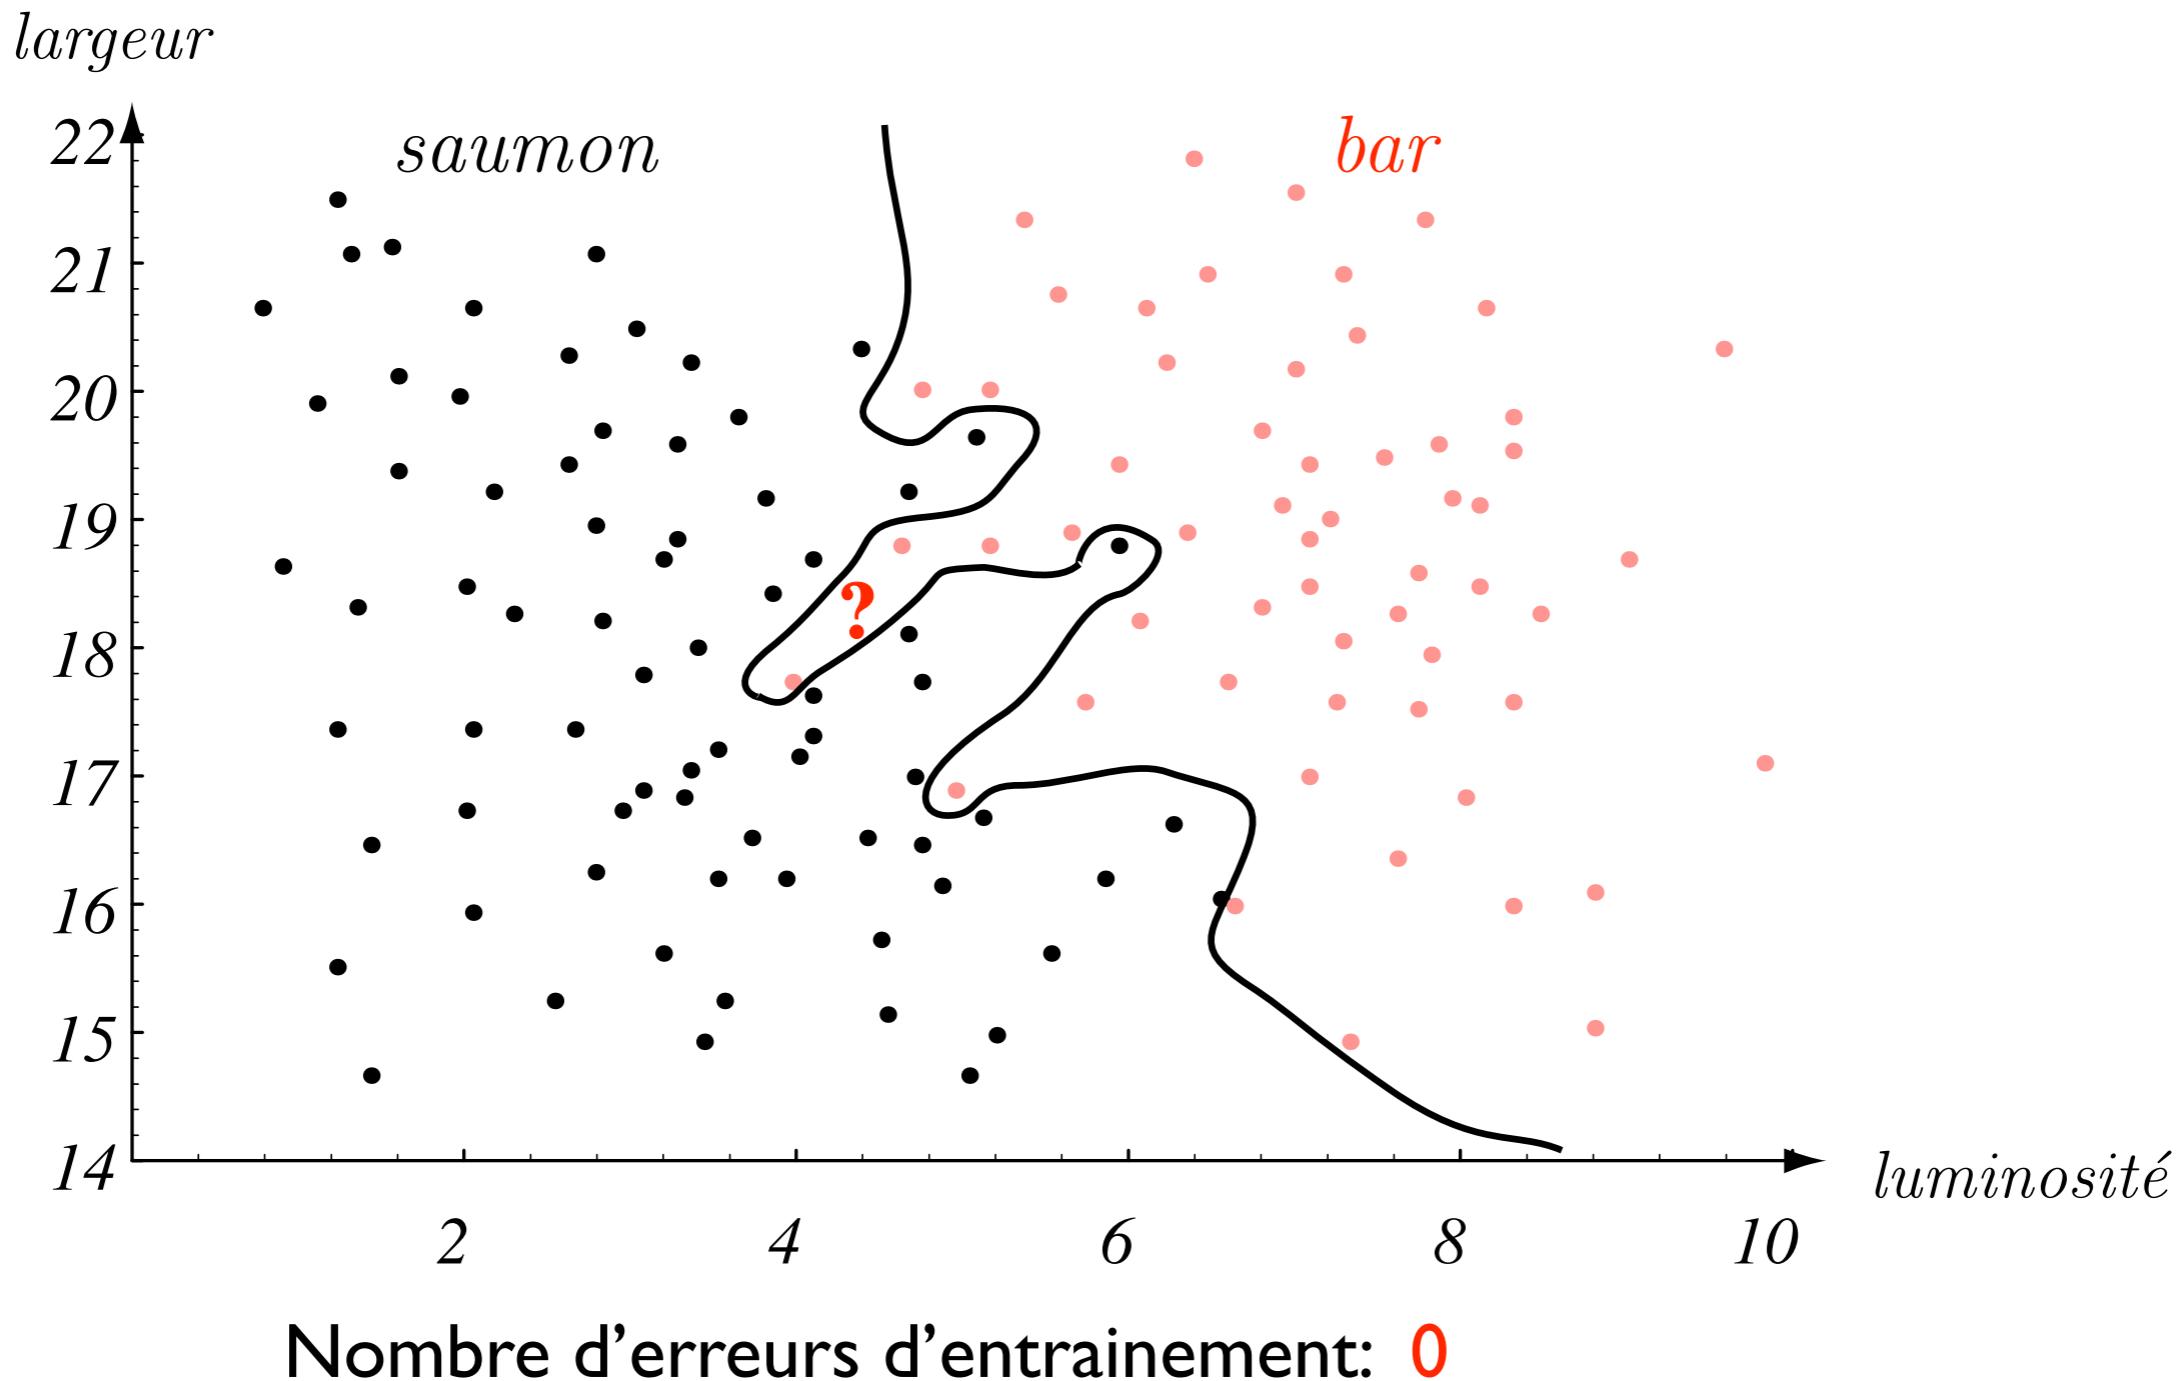
Ex: 2D classification

Linear classifier


- Function family too poor
(too inflexible)
- = **Capacity too low** for this problem
(relative to number of examples)
- => Under-fitting



- Function family too rich
(too flexible)
- = **Capacity too high** for this problem
(relative to the number of examples)
- => Over-fitting

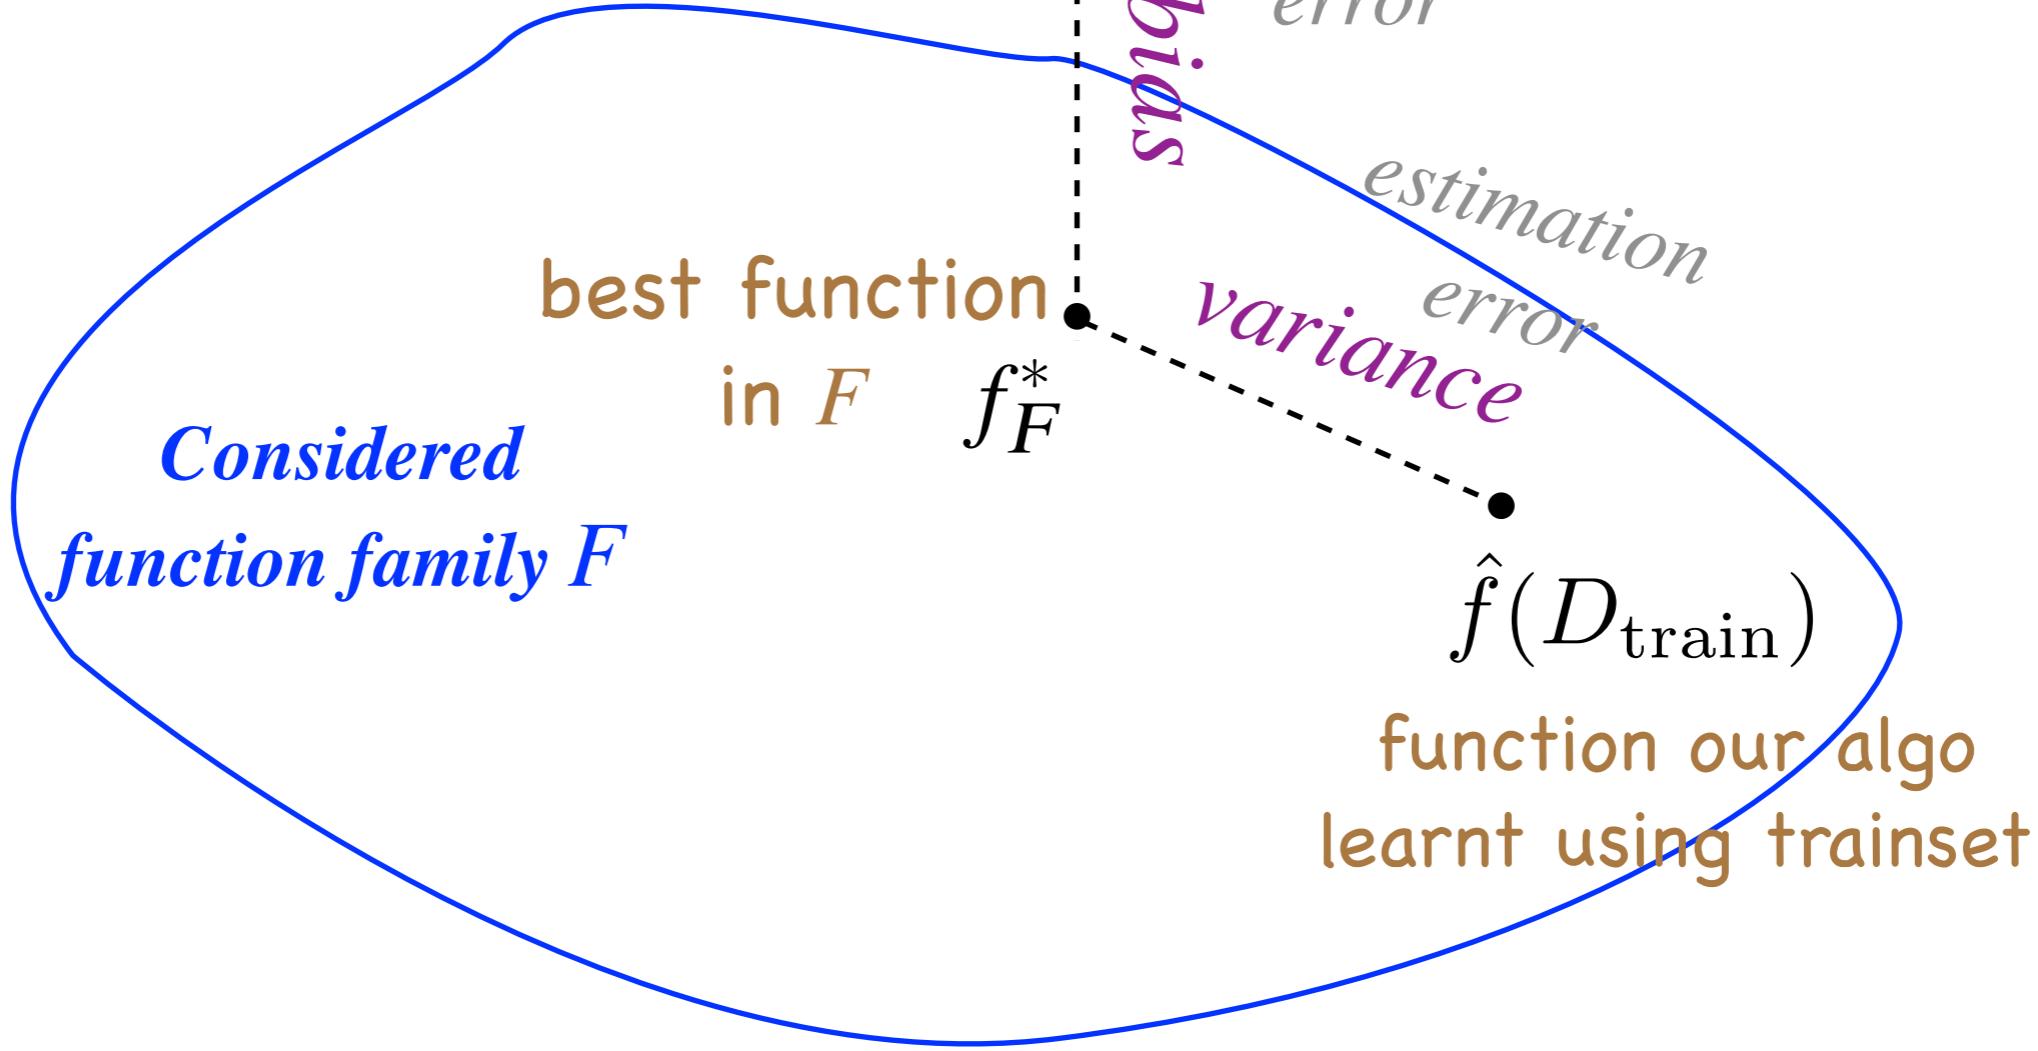


- **Optimal capacity** for this problem
(par rapport à la quantité de données)
- => Best generalization
(on future test points)



Decomposing the generalization error

*Set of all possible
functions
in the universe*

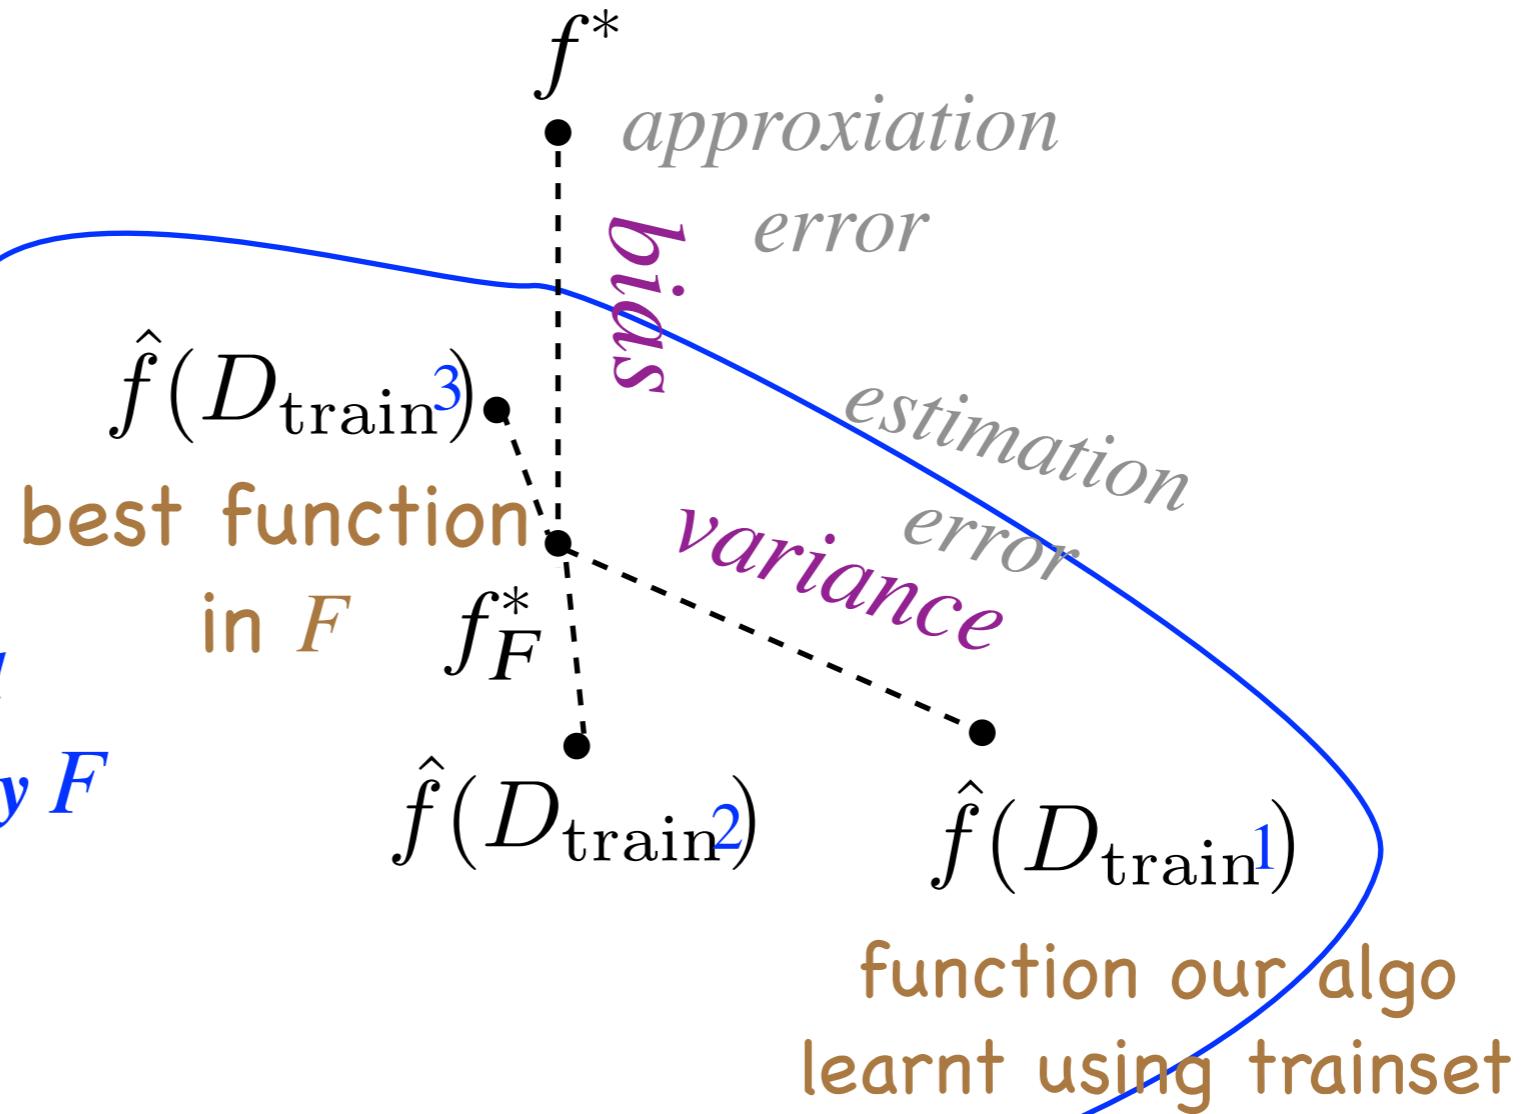


What is responsible for the variance?

*Set of all possible
functions
in the universe*

*Considered
function family F*

best possible
function



Optimal capacity & the bias-variance dilemma

- Choosing richer F : capacity \uparrow
 ⇒ bias \downarrow but variance \uparrow .
- Choosing smaller F : capacity \downarrow
 ⇒ variance \downarrow but bias \uparrow .
- Optimal compromise... will depend on number of examples n
- Bigger n ⇒ variance \downarrow
 So we can afford to increase capacity (to lower the bias)
 ⇒ can use more expressive models
- The best regularizer is more data!

Model selection how to

$D =$

(x_1, y_1)

(x_2, y_2)

— — —

⋮

— — —

(x_N, y_N)

Training set
 D_{train}

Validation set
 D_{valid}

Test set
 D_{test}

Make sure examples are in random order

Split data D in 3: D_{train} D_{valid} D_{test}

Model selection meta-algorithm:

For each considered model (ML algo) A :

For each considered hyper-parameter config λ :

- train model A with hyperparams λ on D_{train}

$$\hat{f}_{A_\lambda} = A_\lambda(D_{\text{train}})$$

- evaluate resulting predictor on D_{valid}
(with preferred evaluation metric)

$$e_{A_\lambda} = \hat{R}(\hat{f}_{A_\lambda}, D_{\text{valid}})$$

Locate A^*, λ^* that yielded best e_{A_λ}

Either return $f^* = f_{A_{\lambda^*}}$

Or retrain and return

$$f^* = A_{\lambda^*}^*(D_{\text{train}} \cup D_{\text{valid}})$$


Finally: compute unbiased estimate of generalization performance of f^* using D_{test}

$$\hat{R}(f^*, D_{\text{test}})$$

D_{test} must never have been used during training or model selection to select, learn, or tune anything.

Ex of model hyper-parameter selection

- Training set error
- Validation set error



Hyper-parameter value which yields smallest error on validation set is 5
(it was 1 for the training set)

Question

What if we selected capacity-control hyper-parameters that yield best performance on the training set?

What would we tend to select?

Is it a good idea? Why?

Model selection procedure summary:

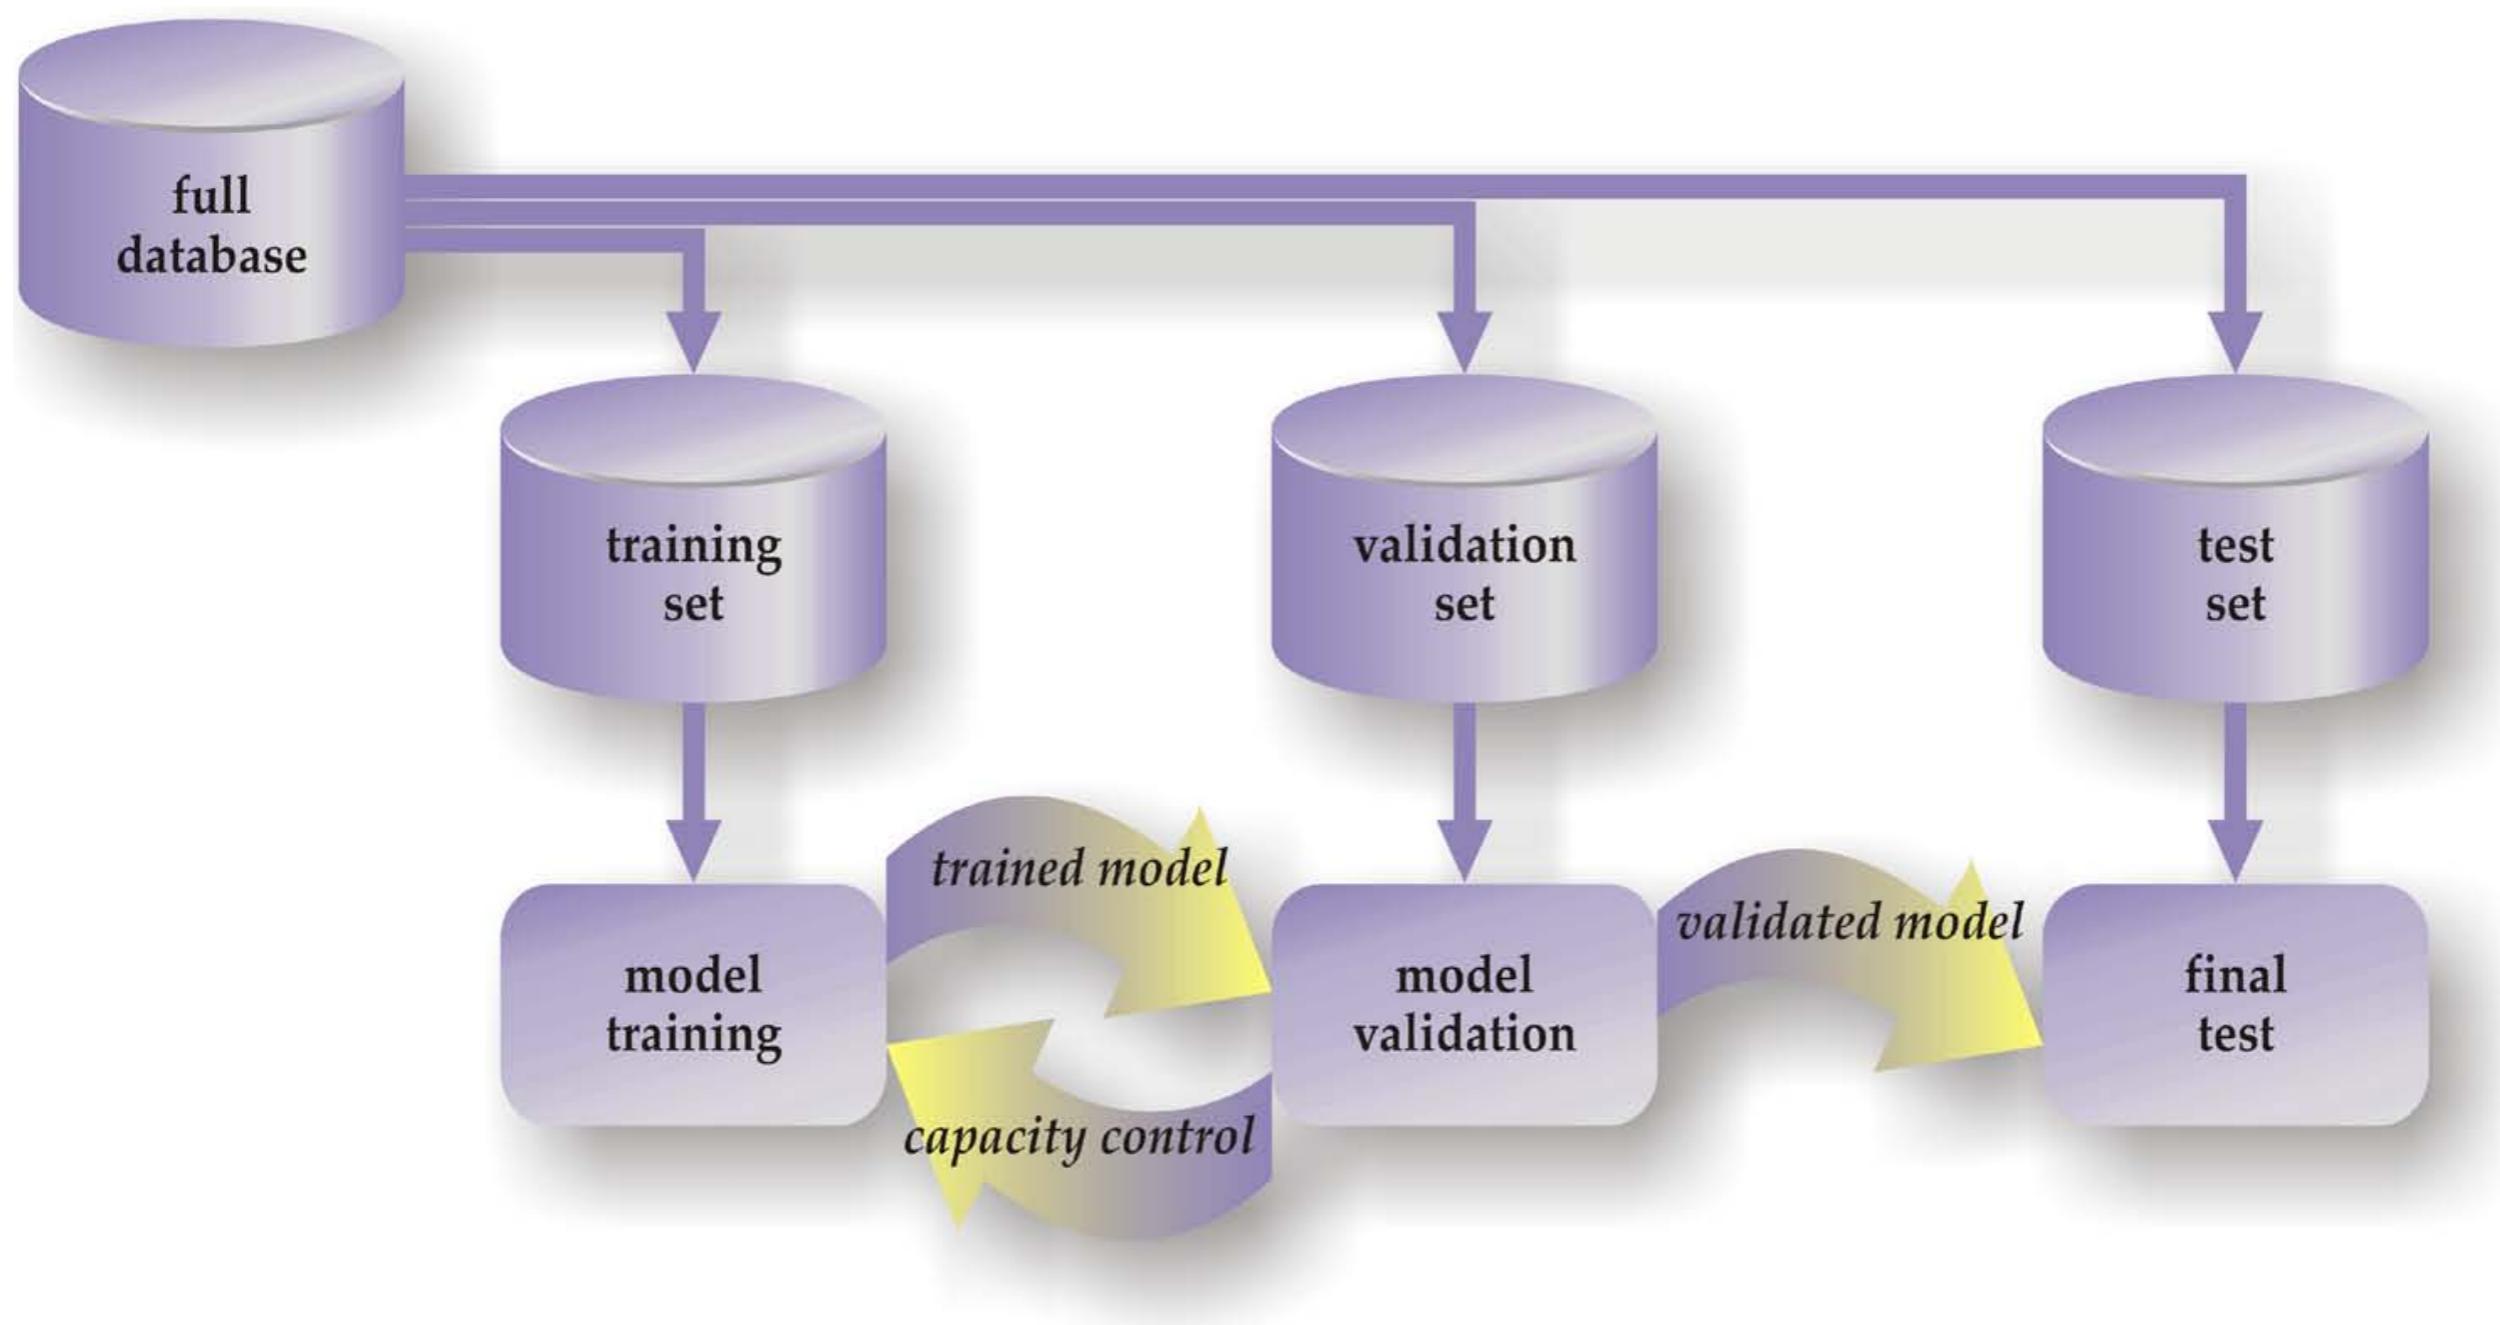
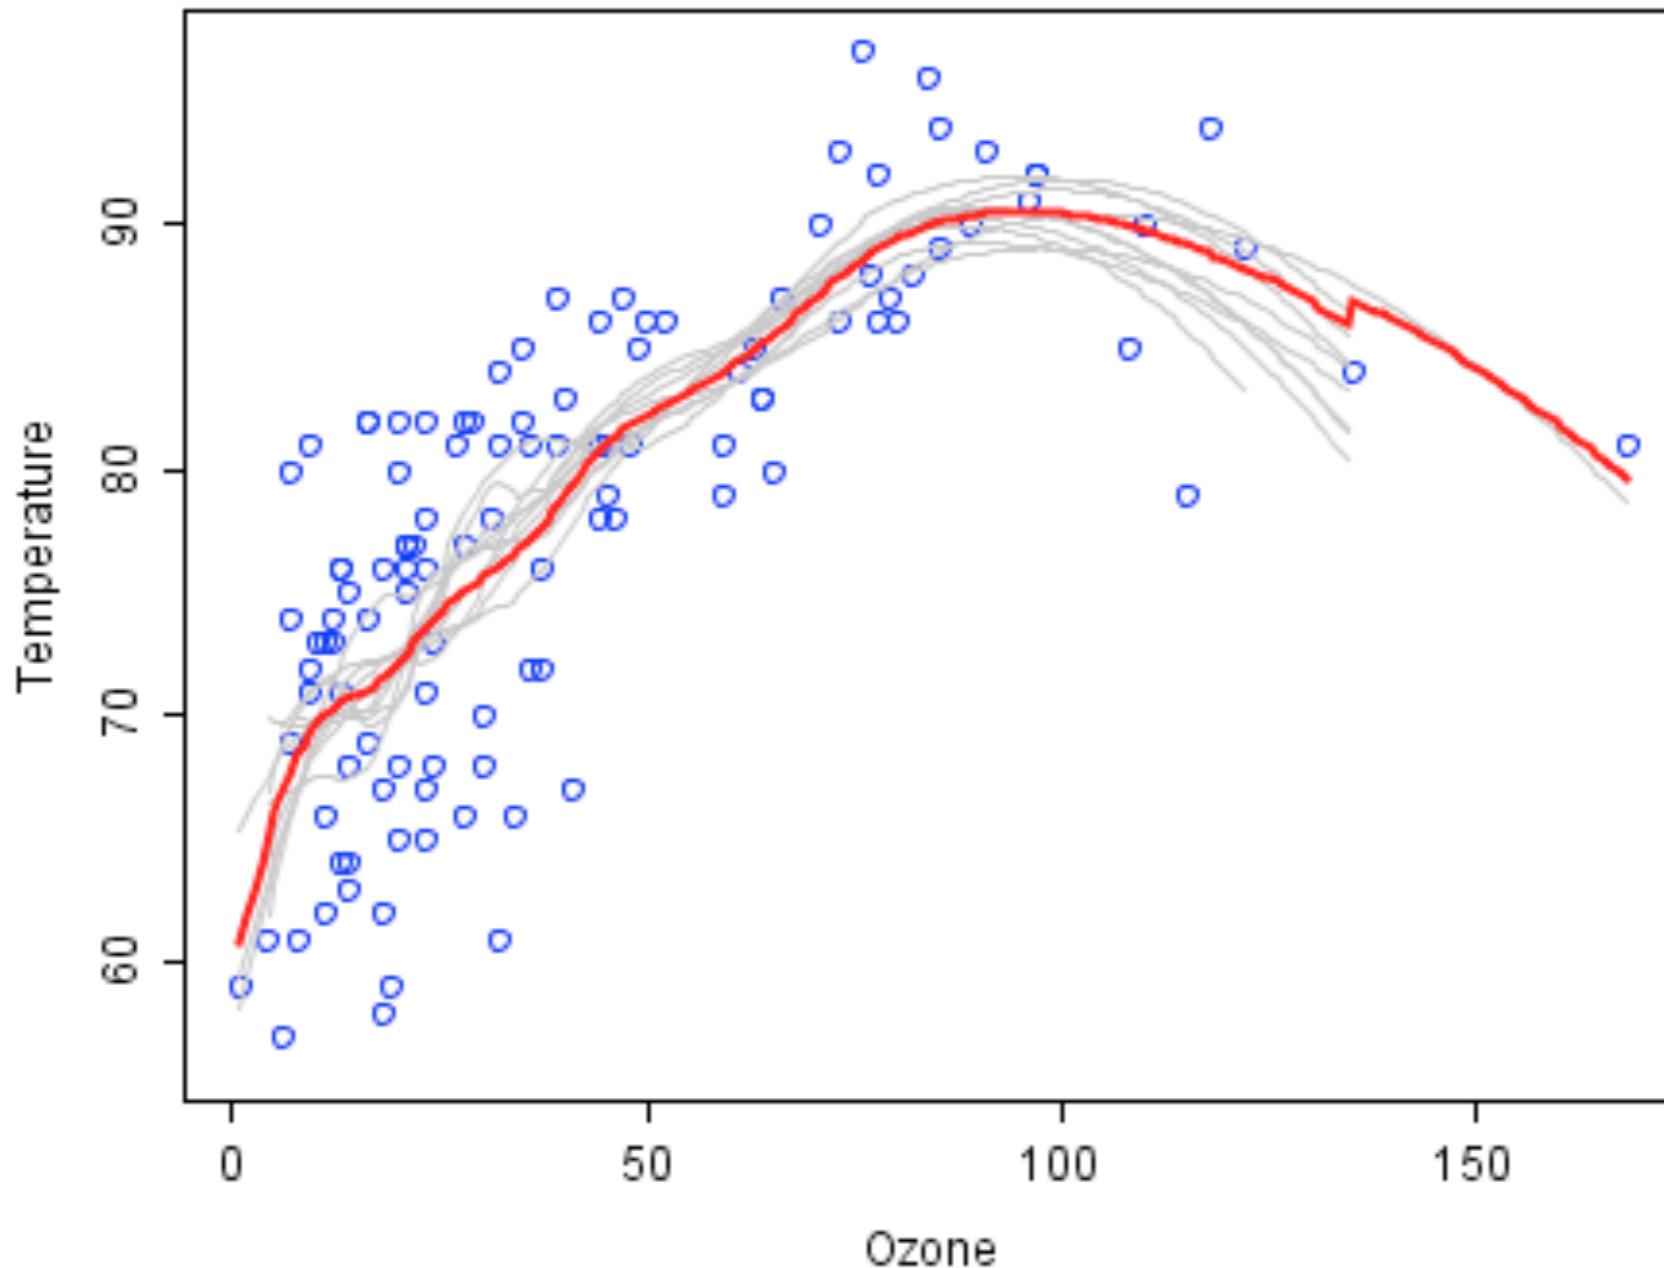


Figure by Nicolas Chapados

Ensemble methods

- Principle: train and combine multiple predictors to good effect
- Bagging: average many high-variance predictors
 - variance ↓
(e.g.: average deep trees → Random decision forests)
- Boosting: build weighted combination of low-capacity classifiers
 - bias ↓ and capacity ↑
(e.g. boosting shallow trees; or linear classifiers)

Bagging for reducing variance on a regression problem



How to obtain non-linear predictor with a linear predictor

Three ways to map x to a feature representation $\tilde{x} = \phi(x)$

- Use an **explicit fixed mapping** (ex: hand-crafted features)
- Use an **implicit fixed mapping**
 - ➡ Kernel Methods (SVMs, Kernel Logistic Regression ...)
- **Learn a parameterized mapping**
 - (i.e. let the ML algo learn the new representation)
 - ➡ **Multilayer feed-forward Neural Networks**
such as **Multilayer Perceptrons (MLP)**

Levels of representation



very high level representation:

CAT

JUMPING

... etc ...

slightly higher level representation

raw input vector representation:

$$\mathcal{X} = \boxed{23 \quad 19 \quad 20} \quad \cdots \quad \boxed{\quad \quad 18}$$

x_1 x_2 x_3 \cdots x_n





Questions ?